

# *Artificial Intelligence: Modeling Human Intelligence with Networks*

---

Jeová Farias Sales Rocha Neto  
jeova\_farias@brown.edu

# Some news

---

- Last Friday presentations.

## Some news

---

- Last Friday presentations.
- Next Friday presentations.

## From the previous chapters

---

- We've learned a powerful optimization technique last class called

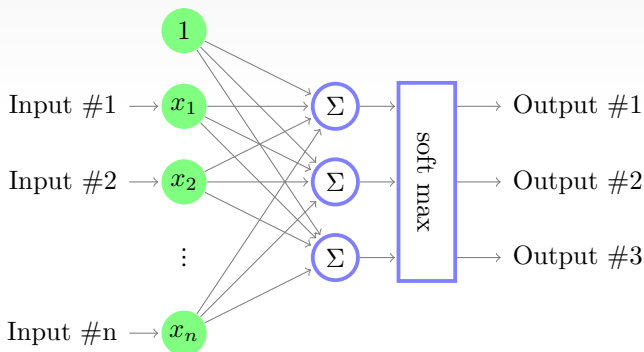
## From the previous chapters

---

- We've learned a powerful optimization technique last class called **Gradient Descent**.

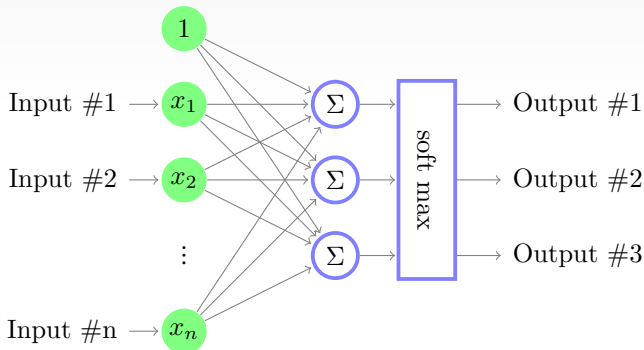
## From the previous chapters

- We've learned a powerful optimization technique last class called **Gradient Descent**.
- We can come back to where we left our previous neural network:



## From the previous chapters

- We've learned a powerful optimization technique last class called **Gradient Descent**.
- We can come back to where we left our previous neural network:



$$\text{Output} = o = [o_1, o_2, o_3] = \text{softmax}(Wx)$$

# Loss Function

---

- What do we wanna minimize with Gradient Descent?



# Loss Function

---

- What do we wanna minimize with Gradient Descent? A loss function (cf. this slide's title)!
- What is the loss we've been working with?

# Loss Function

---

- What do we wanna minimize with Gradient Descent? A loss function (cf. this slide's title)!
- What is the loss we've been working with?
- The zero-one loss (for the perceptron algorithm):

$$\mathcal{L}(X, w) = \sum_{i=1}^m [f_{act}(w \cdot x^i) \neq class_i]$$

# Loss Function

- What do we wanna minimize with Gradient Descent? A loss function (cf. this slide's title)!
- What is the loss we've been working with?
- The zero-one loss (for the perceptron algorithm):

$$\mathcal{L}(X, w) = \sum_{i=1}^m [f_{act}(w \cdot x^i) \neq class_i]$$

- What kind of output is the output of the softmax function?

# Loss Function

- What do we wanna minimize with Gradient Descent? A loss function (cf. this slide's title)!
- What is the loss we've been working with?
- The zero-one loss (for the perceptron algorithm):

$$\mathcal{L}(X, w) = \sum_{i=1}^m [f_{act}(w \cdot x^i) \neq class_i]$$

- What kind of output is the output of the softmax function?
- Does it work in the zero-one loss? (say no)

# Loss Function

- What do we wanna minimize with Gradient Descent? A loss function (cf. this slide's title)!
- What is the loss we've been working with?
- The zero-one loss (for the perceptron algorithm):

$$\mathcal{L}(X, w) = \sum_{i=1}^m [f_{act}(w \cdot x^i) \neq class_i]$$

- What kind of output is the output of the softmax function?
- Does it work in the zero-one loss? (say no) Why?

# Loss Function

- What do we wanna minimize with Gradient Descent? A loss function (cf. this slide's title)!
- What is the loss we've been working with?
- The zero-one loss (for the perceptron algorithm):

$$\mathcal{L}(X, w) = \sum_{i=1}^m [f_{act}(w \cdot x^i) \neq class_i]$$

- What kind of output is the output of the softmax function?
- Does it work in the zero-one loss? (say no) Why?
- What to do? Cry?

# Cross entropy

---

- First with need to define the Cross Entropy between two probability distributions (percentage)  $p, q \in \mathbb{R}^n$  is:

$$\text{CrossEntropy}(p, q) = - \sum_{j=1}^n p_j \log(q_j)$$

# Cross entropy

---

- First with need to define the Cross Entropy between two probability distributions (percentage)  $p, q \in \mathbb{R}^n$  is:

$$\text{CrossEntropy}(p, q) = - \sum_{j=1}^n p_i \log(q_i)$$

Ex1.:  $p = [.5, .2, .3]$ ,  $q = [.3, .3, .4]$ ,

$$\text{CrossEntropy}(p, q) \approx -.5 \times (-1.21) - .2 \times (-1.21) - .3 \times (-0.92) = 1.12$$



## Cross entropy

- First with need to define the Cross Entropy between two probability distributions (percentage)  $p, q \in \mathbb{R}^n$  is:

$$\text{CrossEntropy}(p, q) = - \sum_{j=1}^n p_j \log(q_j)$$

Ex1.:  $p = [.5, .2, .3]$ ,  $q = [.3, .3, .4]$ ,

$$\text{CrossEntropy}(p, q) \approx -.5 \times (-1.21) - .2 \times (-1.21) - .3 \times (-0.92) = 1.12$$

Ex2.:  $p = [1, 0, 0]$ ,  $q = [1, 0, 0]$ ,

$$\text{CrossEntropy}(p, q) = 1 \times (\log(1)) + 0 \times (\log(0)) + 0 \times (\log(0))$$

## Cross entropy

- First with need to define the Cross Entropy between two probability distributions (percentage)  $p, q \in \mathbb{R}^n$  is:

$$\text{CrossEntropy}(p, q) = - \sum_{j=1}^n p_j \log(q_j)$$

Ex1.:  $p = [.5, .2, .3]$ ,  $q = [.3, .3, .4]$ ,

$$\text{CrossEntropy}(p, q) \approx -.5 \times (-1.21) - .2 \times (-1.21) - .3 \times (-0.92) = 1.12$$

Ex2.:  $p = [1, 0, 0]$ ,  $q = [1, 0, 0]$ ,

$$\text{CrossEntropy}(p, q) = 1 \times (\log(1)) + 0 \times (\log(0)) + 0 \times (\log(0)) = 0$$

# Cross entropy

- First with need to define the Cross Entropy between two probability distributions (percentage)  $p, q \in \mathbb{R}^n$  is:

$$\text{CrossEntropy}(p, q) = - \sum_{j=1}^n p_j \log(q_j)$$

Ex1.:  $p = [.5, .2, .3]$ ,  $q = [.3, .3, .4]$ ,

$$\text{CrossEntropy}(p, q) \approx -.5 \times (-1.21) - .2 \times (-1.21) - .3 \times (-0.92) = 1.12$$

Ex2.:  $p = [1, 0, 0]$ ,  $q = [1, 0, 0]$ ,

$$\text{CrossEntropy}(p, q) = 1 \times (\log(1)) + 0 \times (\log(0)) + 0 \times (\log(0)) = 0$$

Ex3.:  $p = [1, 0, 0]$ ,  $q = [0, 1, 0]$ ,

$$\text{CrossEntropy}(p, q) = 1 \times (\log(0)) + 0 \times (\log(1)) + 0 \times (\log(0))$$

# Cross entropy

- First with need to define the Cross Entropy between two probability distributions (percentage)  $p, q \in \mathbb{R}^n$  is:

$$\text{CrossEntropy}(p, q) = - \sum_{j=1}^n p_i \log(q_i)$$

Ex1.:  $p = [.5, .2, .3], q = [.3, .3, .4],$

$$\text{CrossEntropy}(p, q) \approx -.5 \times (-1.21) - .2 \times (-1.21) - .3 \times (-0.92) = 1.12$$

Ex2.:  $p = [1, 0, 0], q = [1, 0, 0],$

$$\text{CrossEntropy}(p, q) = 1 \times (\log(1)) + 0 \times (\log(0)) + 0 \times (\log(0)) = 0$$

Ex3.:  $p = [1, 0, 0], q = [0, 1, 0],$

$$\text{CrossEntropy}(p, q) = 1 \times (\log(0)) + 0 \times (\log(1)) + 0 \times (\log(0)) = \infty$$

- Does the definition of Cross Entropy remind you of something we saw in class?

# Cross entropy

- First with need to define the Cross Entropy between two probability distributions (percentage)  $p, q \in \mathbb{R}^n$  is:

$$\text{CrossEntropy}(p, q) = - \sum_{j=1}^n p_i \log(q_i)$$

Ex1.:  $p = [.5, .2, .3], q = [.3, .3, .4],$

$$\text{CrossEntropy}(p, q) \approx -.5 \times (-1.21) - .2 \times (-1.21) - .3 \times (-0.92) = 1.12$$

Ex2.:  $p = [1, 0, 0], q = [1, 0, 0],$

$$\text{CrossEntropy}(p, q) = 1 \times (\log(1)) + 0 \times (\log(0)) + 0 \times (\log(0)) = 0$$

Ex3.:  $p = [1, 0, 0], q = [0, 1, 0],$

$$\text{CrossEntropy}(p, q) = 1 \times (\log(0)) + 0 \times (\log(1)) + 0 \times (\log(0)) = \infty$$

- Does the definition of Cross Entropy remind you of something we saw in class? Dot product!

# Cross entropy

- First with need to define the Cross Entropy between two probability distributions (percentage)  $p, q \in \mathbb{R}^n$  is:

$$\text{CrossEntropy}(p, q) = - \sum_{j=1}^n p_i \log(q_i)$$

Ex1.:  $p = [.5, .2, .3], q = [.3, .3, .4],$

$$\text{CrossEntropy}(p, q) \approx -.5 \times (-1.21) - .2 \times (-1.21) - .3 \times (-0.92) = 1.12$$

Ex2.:  $p = [1, 0, 0], q = [1, 0, 0],$

$$\text{CrossEntropy}(p, q) = 1 \times (\log(1)) + 0 \times (\log(0)) + 0 \times (\log(0)) = 0$$

Ex3.:  $p = [1, 0, 0], q = [0, 1, 0],$

$$\text{CrossEntropy}(p, q) = 1 \times (\log(0)) + 0 \times (\log(1)) + 0 \times (\log(0)) = \infty$$

- Does the definition of Cross Entropy remind you of something we saw in class? Dot product!

$$\text{CrossEntropy}(p, q) = -p \cdot \log(q)$$

## Cross entropy loss

---

- For the kinds of output of softmax (percentages), there's another loss called **Cross Entropy Loss**.

$$\begin{aligned}\mathcal{L}(X, W) &= \sum_{i=1}^m \text{CrossEntropy}(y^i, o^i) \\ &= \sum_{i=1}^m \text{CrossEntropy}(y^i, \text{softmax}(Wx^i))\end{aligned}$$

## Cross entropy loss

---

- For the kinds of output of softmax (percentages), there's another loss called **Cross Entropy Loss**.

$$\begin{aligned}\mathcal{L}(X, W) &= \sum_{i=1}^m \text{CrossEntropy}(y^i, o^i) \\ &= \sum_{i=1}^m \text{CrossEntropy}(y^i, \text{softmax}(Wx^i))\end{aligned}$$

- There is a detail on  $y_i$  now:
  - ▶ Previously it was either  $-1$  or  $1$ .



# Cross entropy loss

- For the kinds of output of softmax (percentages), there's another loss called **Cross Entropy Loss**.

$$\begin{aligned}\mathcal{L}(X, W) &= \sum_{i=1}^m \text{CrossEntropy}(y^i, o^i) \\ &= \sum_{i=1}^m \text{CrossEntropy}(y^i, \text{softmax}(Wx^i))\end{aligned}$$

- There is a detail on  $y_i$  now:
  - ▶ Previously it was either  $-1$  or  $1$ .
  - ▶ Now we know that the classes are  $0, 1$  or  $2$ .

# Cross entropy loss

- For the kinds of output of softmax (percentages), there's another loss called **Cross Entropy Loss**.

$$\begin{aligned}\mathcal{L}(X, W) &= \sum_{i=1}^m \text{CrossEntropy}(y^i, o^i) \\ &= \sum_{i=1}^m \text{CrossEntropy}(y^i, \text{softmax}(Wx^i))\end{aligned}$$

- There is a detail on  $y_i$  now:
  - ▶ Previously it was either  $-1$  or  $1$ .
  - ▶ Now we know that the classes are  $0, 1$  or  $2$ .
  - ▶ So, we should create, for each  $y_i$ , a vector  $Y^i$  *one-hot*.

# Cross entropy loss

- For the kinds of output of softmax (percentages), there's another loss called **Cross Entropy Loss**.

$$\begin{aligned}\mathcal{L}(X, W) &= \sum_{i=1}^m \text{CrossEntropy}(y^i, o^i) \\ &= \sum_{i=1}^m \text{CrossEntropy}(y^i, \text{softmax}(Wx^i))\end{aligned}$$

- There is a detail on  $y_i$  now:
  - ▶ Previously it was either  $-1$  or  $1$ .
  - ▶ Now we know that the classes are  $0, 1$  or  $2$ .
  - ▶ So, we should create, for each  $y_i$ , a vector  $Y^i$  *one-hot*.
  - ▶ Example: if  $y_i = 1$ ,  $Y^i = [0, 1, 0]$ . If  $y_i = 0$ ,  $Y^i = [1, 0, 0]$ .

# Cross entropy loss

- For the kinds of output of softmax (percentages), there's another loss called **Cross Entropy Loss**.

$$\begin{aligned}\mathcal{L}(X, W) &= \sum_{i=1}^m \text{CrossEntropy}(y^i, o^i) \\ &= \sum_{i=1}^m \text{CrossEntropy}(y^i, \text{softmax}(Wx^i))\end{aligned}$$

- There is a detail on  $y_i$  now:
  - ▶ Previously it was either  $-1$  or  $1$ .
  - ▶ Now we know that the classes are  $0, 1$  or  $2$ .
  - ▶ So, we should create, for each  $y_i$ , a vector  $Y^i$  *one-hot*.
  - ▶ Example: if  $y_i = 1$ ,  $Y^i = [0, 1, 0]$ . If  $y_i = 0$ ,  $Y^i = [1, 0, 0]$ .
- Jupyter Notebooks!

# Gradient Descent on the Multilayer perceptron

- Ok, now we need to find a  $W = [w^0, w^1, w^2]$  that minimizes  $\mathcal{L}(X, W)$ .

$$\mathcal{L}(X, W) = \sum_{i=1}^m \text{CrossEntropy}(y^i, \text{softmax}(W x^i)) \quad (1)$$

# Gradient Descent on the Multilayer perceptron

- Ok, now we need to find a  $W = [w^0, w^1, w^2]$  that minimizes  $\mathcal{L}(X, W)$ .

$$\mathcal{L}(X, W) = \sum_{i=1}^m \text{CrossEntropy}(y^i, \text{softmax}(Wx^i)) \quad (1)$$

- Let's use Gradient Descent!

# Gradient Descent on the Multilayer perceptron

- Ok, now we need to find a  $W = [w^0, w^1, w^2]$  that minimizes  $\mathcal{L}(X, W)$ .

$$\mathcal{L}(X, W) = \sum_{i=1}^m \text{CrossEntropy}(y^i, \text{softmax}(W x^i)) \quad (1)$$

- Let's use Gradient Descent! Wait,

# Gradient Descent on the Multilayer perceptron

- Ok, now we need to find a  $W = [w^0, w^1, w^2]$  that minimizes  $\mathcal{L}(X, W)$ .

$$\mathcal{L}(X, W) = \sum_{i=1}^m \text{CrossEntropy}(y^i, \text{softmax}(Wx^i)) \quad (1)$$

- Let's use Gradient Descent! Wait, how do we do it?



# Gradient Descent on the Multilayer perceptron

- Ok, now we need to find a  $W = [w^0, w^1, w^2]$  that minimizes  $\mathcal{L}(X, W)$ .

$$\mathcal{L}(X, W) = \sum_{i=1}^m \text{CrossEntropy}(y^i, \text{softmax}(Wx^i)) \quad (1)$$

- Let's use Gradient Descent! Wait, how do we do it?
- First notice that we want to find  $W$  that minimizes  $\mathcal{L}(X, W)$ .

# Gradient Descent on the Multilayer perceptron

- Ok, now we need to find a  $W = [w^0, w^1, w^2]$  that minimizes  $\mathcal{L}(X, W)$ .

$$\mathcal{L}(X, W) = \sum_{i=1}^m \text{CrossEntropy}(y^i, \text{softmax}(W x^i)) \quad (1)$$

- Let's use Gradient Descent! Wait, how do we do it?
- First notice that we want to find  $W$  that minimizes  $\mathcal{L}(X, W)$ .
- So the gradient descent iterations are like:

$$W = W - \eta \frac{d\mathcal{L}(X, W)}{dW},$$

where  $\eta$  is the step size.

# Gradient Descent on the Multilayer perceptron

- Ok, now we need to find a  $W = [w^0, w^1, w^2]$  that minimizes  $\mathcal{L}(X, W)$ .

$$\mathcal{L}(X, W) = \sum_{i=1}^m \text{CrossEntropy}(y^i, \text{softmax}(W x^i)) \quad (1)$$

- Let's use Gradient Descent! Wait, how do we do it?
- First notice that we want to find  $W$  that minimizes  $\mathcal{L}(X, W)$ .
- So the gradient descent iterations are like:

$$W = W - \eta \frac{d\mathcal{L}(X, W)}{dW},$$

where  $\eta$  is the step size.

- How to compute the  $\frac{d\mathcal{L}(X, W)}{dW}$ ?

# Gradient Descent on the Multilayer perceptron

- Chain Rule!

# Gradient Descent on the Multilayer perceptron

- Chain Rule!
- First, write  $\mathcal{L}(X, W)$  as  $\mathcal{L}(X, W) = \sum_{i=1}^m h_i(W)$  and  $h_i(W) = \text{CrossEntropy}(y^i, \text{softmax}(Wx^i))$

# Gradient Descent on the Multilayer perceptron

- Chain Rule!
- First, write  $\mathcal{L}(X, W)$  as  $\mathcal{L}(X, W) = \sum_{i=1}^m h_i(W)$  and  $h_i(W) = \text{CrossEntropy}(y^i, \text{softmax}(Wx^i))$
- Now we have that:

$$\begin{aligned} h_i(W) &= g_i(f_i(W)) \text{ with} \\ g_i(v) &= \text{CrossEntropy}(y^i, \text{softmax}(v)) \\ f(W) &= Wx^i \end{aligned} \tag{2}$$

# Gradient Descent on the Multilayer perceptron

- Chain Rule!
- First, write  $\mathcal{L}(X, W)$  as  $\mathcal{L}(X, W) = \sum_{i=1}^m h_i(W)$  and  $h_i(W) = \text{CrossEntropy}(y^i, \text{softmax}(Wx^i))$
- Now we have that:

$$\begin{aligned}
 h_i(W) &= g_i(f_i(W)) \text{ with} & (2) \\
 g_i(v) &= \text{CrossEntropy}(y^i, \text{softmax}(v)) \\
 f(W) &= Wx^i
 \end{aligned}$$

- Then, we have:

$$\frac{dh_i(W)}{dW} = \frac{dg(f(W))}{df(W)} \frac{df(W)}{dW}$$

# Gradient Descent on the Multilayer perceptron

- Chain Rule!
- First, write  $\mathcal{L}(X, W)$  as  $\mathcal{L}(X, W) = \sum_{i=1}^m h_i(W)$  and  $h_i(W) = \text{CrossEntropy}(y^i, \text{softmax}(Wx^i))$
- Now we have that:

$$\begin{aligned}
 h_i(W) &= g_i(f_i(W)) \text{ with} \\
 g_i(v) &= \text{CrossEntropy}(y^i, \text{softmax}(v)) \\
 f(W) &= Wx^i
 \end{aligned} \tag{2}$$

- Then, we have:

$$\frac{dh_i(W)}{dW} = \frac{dg(f(W))}{df(W)} \frac{df(W)}{dW}$$

- Notice that (after a looot of algebra):

$$\frac{dg(f(W))}{df(W)} = \text{softmax}(Wx^i) - y^i$$



# Gradient Descent on the Multilayer perceptron

- Chain Rule!
- First, write  $\mathcal{L}(X, W)$  as  $\mathcal{L}(X, W) = \sum_{i=1}^m h_i(W)$  and  $h_i(W) = \text{CrossEntropy}(y^i, \text{softmax}(Wx^i))$
- Now we have that:

$$\begin{aligned}
 h_i(W) &= g_i(f_i(W)) \text{ with} \\
 g_i(v) &= \text{CrossEntropy}(y^i, \text{softmax}(v)) \\
 f(W) &= Wx^i
 \end{aligned} \tag{2}$$

- Then, we have:

$$\frac{dh_i(W)}{dW} = \frac{dg(f(W))}{df(W)} \frac{df(W)}{dW}$$

- Notice that (after a looot of algebra):

$$\frac{dg(f(W))}{df(W)} = \text{softmax}(Wx^i) - y^i = o^i - y^i,$$

# Gradient Descent on the Multilayer perceptron

- Chain Rule!
- First, write  $\mathcal{L}(X, W)$  as  $\mathcal{L}(X, W) = \sum_{i=1}^m h_i(W)$  and  $h_i(W) = \text{CrossEntropy}(y^i, \text{softmax}(Wx^i))$
- Now we have that:

$$\begin{aligned}
 h_i(W) &= g_i(f_i(W)) \text{ with} \\
 g_i(v) &= \text{CrossEntropy}(y^i, \text{softmax}(v)) \\
 f(W) &= Wx^i
 \end{aligned} \tag{2}$$

- Then, we have:

$$\frac{dh_i(W)}{dW} = \frac{dg(f(W))}{df(W)} \frac{df(W)}{dW}$$

- Notice that (after a looot of algebra):

$$\frac{dg(f(W))}{df(W)} = \text{softmax}(Wx^i) - y^i = o^i - y^i, \quad \frac{df(W)}{dW} = x^i$$

# Gradient Descent on the Multilayer perceptron

- And finally:

$$\begin{aligned}
 \frac{d\mathcal{L}(X, W)}{dW} &= \frac{d(\sum_{i=1}^m h_i(W))}{dW} \\
 &= \sum_{i=1}^m \frac{dh_i(W)}{dW} \\
 &= \sum_{i=1}^m (y^i - \text{softmax}(Wx^i)) \cdot x^i, \\
 &= \sum_{i=1}^m (o^i - y^i) \cdot x^i, \\
 &= (O - Y)X,
 \end{aligned}$$

with  $y_i$  being a one-hot vector and  $Y$  the concatenation of  $y^i$ .

- So the Gradient Descent update rule is

$$W = W - \eta \sum_{i=1}^m (y^i - \text{softmax}(Wx^i))x^i,$$

# Gradient Descent on the Multilayer perceptron

- And finally:

$$\begin{aligned}
 \frac{d\mathcal{L}(X, W)}{dW} &= \frac{d(\sum_{i=1}^m h_i(W))}{dW} \\
 &= \sum_{i=1}^m \frac{dh_i(W)}{dW} \\
 &= \sum_{i=1}^m (y^i - \text{softmax}(Wx^i)) \cdot x^i, \\
 &= \sum_{i=1}^m (o^i - y^i) \cdot x^i, \\
 &= (O - Y)X,
 \end{aligned}$$

with  $y_i$  being a one-hot vector and  $Y$  the concatenation of  $y^i$ .

- So the Gradient Descent update rule is

$$W = W - \eta \sum_{i=1}^m (y^i - \text{softmax}(Wx^i))x^i,$$

# Gradient Descent on the Multilayer perceptron

---

**Algorithm 1** *GradientDescent – MultilayerPerceptron*

---

**Input:** A matrix of points  $X$ , a vector of classes  $y$

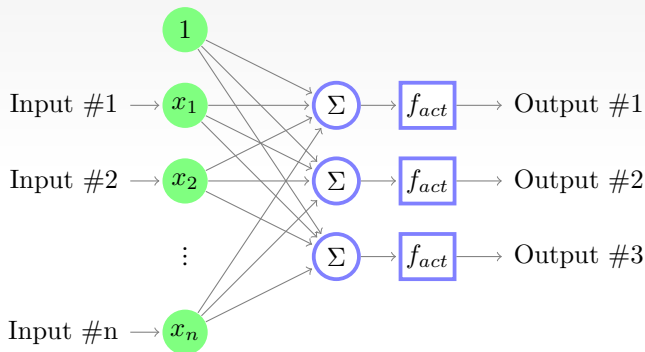
**Output:** A matrix of weights  $W$ .

- 1: Add bias to  $X$
  - 2: Find  $Y$  as the one hot version of  $y$ .
  - 3: Initialize  $W$
  - 4:  $\eta = 1$
  - 5: **for** a given number of epochs (repetitions) **do**
  - 6:    $O = \text{softmax}(WX) \mapsto$  Feed Forward step
  - 7:    $\frac{d\mathcal{L}(X,W)}{dW} = (O - Y)X \mapsto$  Back Propagation step
  - 8:    $W = W - \eta \frac{d\mathcal{L}(X,W)}{dW}$
  - 9: **end for**
-

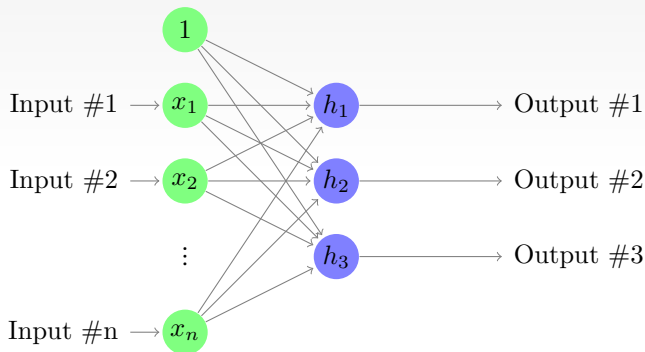
## Neural Evolution (*Part II*)

---

# Neural Evolution




# Neural Evolution

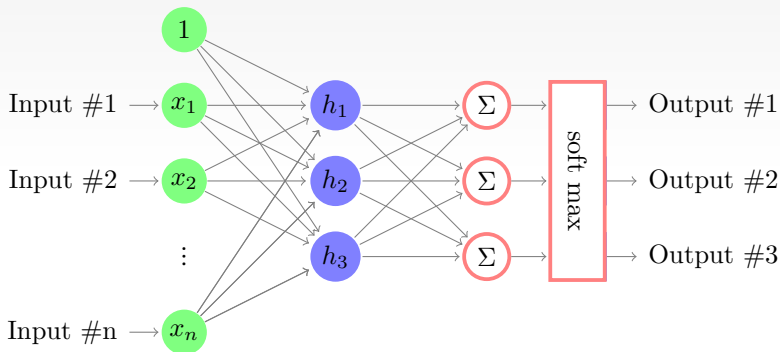




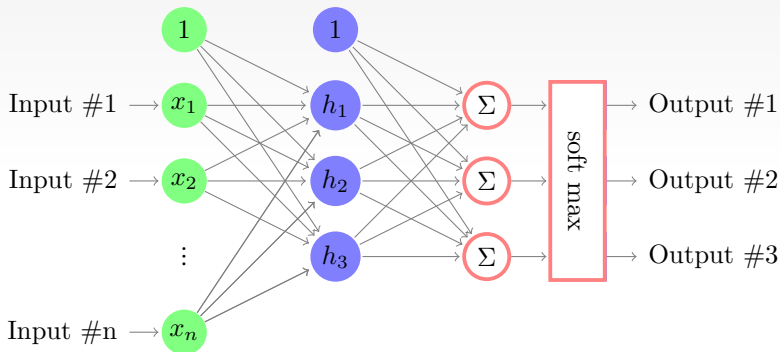
# Neural Evolution



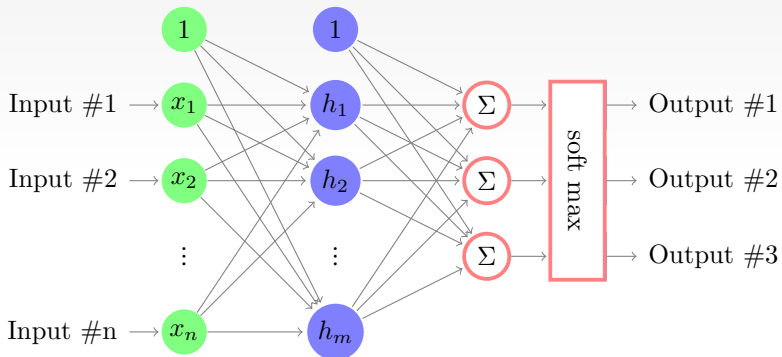
# Neural Evolution



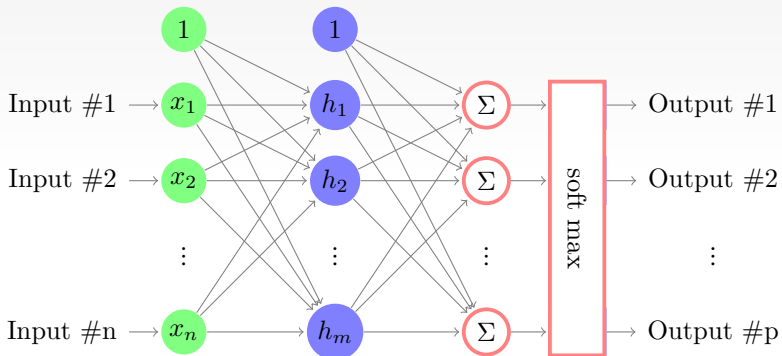
# Neural Evolution




# Neural Evolution



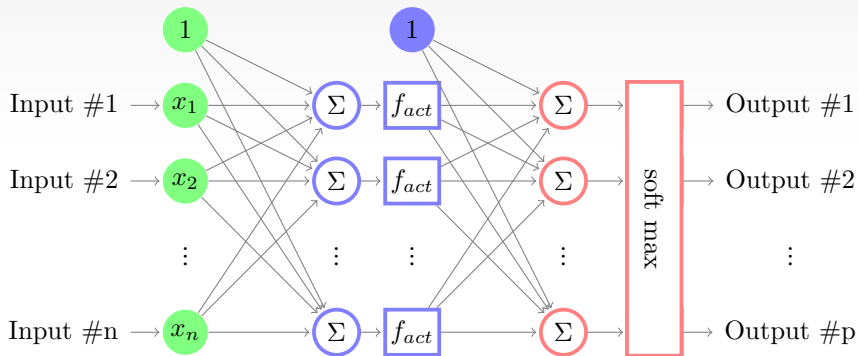
# Neural Evolution



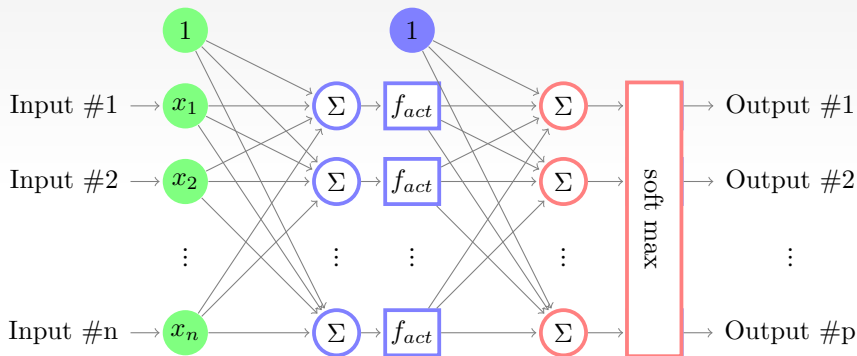
# Neural Evolution



# Neural Evolution




# Neural Evolution




$$o = \text{softmax}(Wh) = \text{softmax}(W(f_{act}(W_h x)))$$



# Activation Functions



(a) ReLU:  $f(x) = \max(0, x)$



(b) tanh:  $f(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$